

Om Chiragbhai Choksi

INTERNSHIP TASK OASIS INFOBYTE

Idea: Customer Segmentation Analysis

Project Description:

The aim of this data analytics project is to perform customer segmentation analysis for an ecommerce company. By analyzing customer behavior and purchase patterns, the goal is to group customers into distinct segments. This segmentation can inform targeted marketing strategies, improve customer satisfaction, and enhance overall business strategies.

Import Libraries

```
!pip install opendatasets
```

```
Collecting opendatasets
```

```
  Downloading opendatasets-0.1.22-py3-none-any.whl.metadata (9.2 kB)
```

```
Requirement already satisfied: tqdm in /usr/local/lib/python3.11/dist-packages (from opendatasets) (4.67.1)
```

```
Requirement already satisfied: kaggle in
```

```
/usr/local/lib/python3.11/dist-packages (from opendatasets) (1.7.4.5)
```

```
Requirement already satisfied: click in
```

```
/usr/local/lib/python3.11/dist-packages (from opendatasets) (8.2.0)
```

```
Requirement already satisfied: bleach in
```

```
/usr/local/lib/python3.11/dist-packages (from kaggle->opendatasets) (6.2.0)
```

```
Requirement already satisfied: certifi>=14.05.14 in
```

```
/usr/local/lib/python3.11/dist-packages (from kaggle->opendatasets) (2025.4.26)
```

```
Requirement already satisfied: charset-normalizer in
```

```
/usr/local/lib/python3.11/dist-packages (from kaggle->opendatasets) (3.4.2)
```

```
Requirement already satisfied: idna in /usr/local/lib/python3.11/dist-packages (from kaggle->opendatasets) (3.10)
```

```
Requirement already satisfied: protobuf in
```

```
/usr/local/lib/python3.11/dist-packages (from kaggle->opendatasets) (5.29.4)
```

```
Requirement already satisfied: python-dateutil>=2.5.3 in
```

```
/usr/local/lib/python3.11/dist-packages (from kaggle->opendatasets) (2.9.0.post0)
```

```
Requirement already satisfied: python-slugify in
```

```
/usr/local/lib/python3.11/dist-packages (from kaggle->opendatasets) (8.0.4)
```

```
Requirement already satisfied: requests in
```

```

/usr/local/lib/python3.11/dist-packages (from kaggle->opendatasets)
(2.32.3)
Requirement already satisfied: setuptools>=21.0.0 in
/usr/local/lib/python3.11/dist-packages (from kaggle->opendatasets)
(75.2.0)
Requirement already satisfied: six>=1.10 in
/usr/local/lib/python3.11/dist-packages (from kaggle->opendatasets)
(1.17.0)
Requirement already satisfied: text-unidecode in
/usr/local/lib/python3.11/dist-packages (from kaggle->opendatasets)
(1.3)
Requirement already satisfied: urllib3>=1.15.1 in
/usr/local/lib/python3.11/dist-packages (from kaggle->opendatasets)
(2.4.0)
Requirement already satisfied: webencodings in
/usr/local/lib/python3.11/dist-packages (from kaggle->opendatasets)
(0.5.1)
Downloading opendatasets-0.1.22-py3-none-any.whl (15 kB)
Installing collected packages: opendatasets
Successfully installed opendatasets-0.1.22

```

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
from sklearn.decomposition import PCA
from sklearn.metrics import silhouette_score
import opendatasets as op
from scipy.stats import pointbiserialr

```

Download Dataset using opendatasets by jovian

```

dataset_url = 'https://www.kaggle.com/jackdaoud/marketing-data'
op.download(dataset_url)
#{"username":"omchoksi04","key":"5bd54ee741a1835da715caead453e032"}

```

Please provide your Kaggle credentials to download this dataset. Learn more: <http://bit.ly/kaggle-creds>

Your Kaggle username: omchoksi04

Your Kaggle Key:

Dataset URL: <https://www.kaggle.com/datasets/jackdaoud/marketing-data>

Downloading marketing-data.zip to ./marketing-data

100%|██████████| 643k/643k [00:00<00:00, 645MB/s]

Data

The data contains 2,205 observations and 39 columns. The dataset description on the card does not match the actual columns in the dataset. The below list contains actual columns from the dataset and the assumed descriptions from the column's names.

Feature	Description	Comment
AcceptedCmp1	1 if customer accepted the offer in the 1st campaign, 0 otherwise	
AcceptedCmp2	1 if customer accepted the offer in the 2nd campaign, 0 otherwise	
AcceptedCmp3	1 if customer accepted the offer in the 3rd campaign, 0 otherwise	
AcceptedCmp4	1 if customer accepted the offer in the 4th campaign, 0 otherwise	
AcceptedCmp5	1 if customer accepted the offer in the 5th campaign, 0 otherwise	
AcceptedCmpOverall	overall number of accepted campaigns	This column was added from the list of actual columns
Response	1 if customer accepted the offer in the last campaign, 0 otherwise	
Complain	1 if customer complained in the last 2 years	
DtCustomer	date of customer's enrolment with the company	There is no such column in the dataset
Customer_Days	number of days since registration as a customer	
Education	customer's level of education	There is no such column in the actual dataset
education_2n Cycle	customer has secondary education	This column was added from the list of actual columns
education_Basic	customer has basic education	This column was added from the list of actual columns
education_Graduation	Customer has a bachelor degree	This column was added from the list of actual columns
education_Master	Customer has a masters degree	This column was added from the list of actual

Feature	Description	Comment
education_PhD	Customer has a PhD	columns This column was added from the list of actual columns
Marital	customer's marital status.	There is no such column in the actual dataset
marital_Divorced	1 if customer is divorced, 0 otherwise.	This column was added from the list of actual columns
marital_Married	1 if customer is married, 0 otherwise.	This column was added from the list of actual columns
marital_Single	1 if customer is single, 0 otherwise.	This column was added from the list of actual columns
marital_Together	1 if customer is in relationship, 0 otherwise.	This column was added from the list of actual columns
marital_Widow	1 if customer is a widow / widower, 0 otherwise	
Kidhome	number of small children in customer's household	
Teenhome	number of teenagers in customer's household	
Income	customer's yearly household income	
MntFishProducts	amount spent on fish products in the last 2 years	
MntMeatProducts	amount spent on meat products in the last 2 years	
MntFruits	amount spent on fruits products in the last 2 years	
MntSweetProducts	amount spent on sweet products in the last 2 years	
MntWines	amount spent on wine products in the last 2 years	
MntGoldProds	amount spent on gold products in the last 2 years	
NumDealsPurchases	number of purchases made with discount	
NumCatalogPurchases	number of purchases made using catalogue	
NumStorePurchases	number of purchases made directly in	

Feature	Description	Comment
	stores	
NumWebPurchases	number of purchases made through company's web site	
NumWebVisitsMonth	number of visits to company's web site in the last month	
Recency	number of days since the last purchase	
Z_CostContact		This column was added from the list of actual columns
Z_Revenue		This column was added from the list of actual columns
Age	Age of customer	This column was added from the list of actual columns
MntTotal	Total amount spent on all the products	This column was added from the list of actual columns
MntRegularProds		This column was added from the list of actual columns

```
df=pd.read_csv("/content/marketing-data/ifood_df.csv")
df.head()
```

```
{"type": "dataframe", "variable_name": "df"}
```

```
df.describe(),df.shape
```

```
(
Income      Kidhome      Teenhome      Recency
MntWines \
count      2205.000000      2205.000000      2205.000000      2205.000000
2205.000000
mean       51622.094785         0.442177         0.506576         49.009070
306.164626
std        20713.063826         0.537132         0.544380         28.932111
337.493839
min         1730.000000         0.000000         0.000000         0.000000
0.000000
25%         35196.000000         0.000000         0.000000         24.000000
24.000000
50%         51287.000000         0.000000         0.000000         49.000000
178.000000
75%         68281.000000         1.000000         1.000000         74.000000
507.000000
```

max	113734.000000	2.000000	2.000000	99.000000
-----	---------------	----------	----------	-----------

1493.000000

	MntFruits	MntMeatProducts	MntFishProducts
MntSweetProducts \			
count	2205.000000	2205.000000	2205.000000
2205.000000			
mean	26.403175	165.312018	37.756463
27.128345			
std	39.784484	217.784507	54.824635
41.130468			
min	0.000000	0.000000	0.000000
0.000000			
25%	2.000000	16.000000	3.000000
1.000000			
50%	8.000000	68.000000	12.000000
8.000000			
75%	33.000000	232.000000	50.000000
34.000000			
max	199.000000	1725.000000	259.000000
262.000000			

	MntGoldProds	...	marital_Together	marital_Widow
education_2n Cycle \				
count	2205.000000	...	2205.000000	2205.000000
2205.000000				
mean	44.057143	...	0.257596	0.034467
0.089796				
std	51.736211	...	0.437410	0.182467
0.285954				
min	0.000000	...	0.000000	0.000000
0.000000				
25%	9.000000	...	0.000000	0.000000
0.000000				
50%	25.000000	...	0.000000	0.000000
0.000000				
75%	56.000000	...	1.000000	0.000000
0.000000				
max	321.000000	...	1.000000	1.000000
1.000000				

	education_Basic	education_Graduation	education_Master
education_PhD \			
count	2205.000000	2205.000000	2205.000000
2205.000000			
mean	0.024490	0.504762	0.165079
0.215873			
std	0.154599	0.500091	0.371336
0.411520			

min	0.000000	0.000000	0.000000
0.000000			
25%	0.000000	0.000000	0.000000
0.000000			
50%	0.000000	1.000000	0.000000
0.000000			
75%	0.000000	1.000000	0.000000
0.000000			
max	1.000000	1.000000	1.000000
1.000000			

	MntTotal	MntRegularProds	AcceptedCmpOverall
count	2205.000000	2205.000000	2205.000000
mean	562.764626	518.707483	0.29932
std	575.936911	553.847248	0.68044
min	4.000000	-283.000000	0.00000
25%	56.000000	42.000000	0.00000
50%	343.000000	288.000000	0.00000
75%	964.000000	884.000000	0.00000
max	2491.000000	2458.000000	4.00000

```
[8 rows x 39 columns],
(2205, 39))
```

```
df.columns
```

```
Index(['Income', 'Kidhome', 'Teenhome', 'Recency', 'MntWines',
'MntFruits',
'MntMeatProducts', 'MntFishProducts', 'MntSweetProducts',
'MntGoldProds', 'NumDealsPurchases', 'NumWebPurchases',
'NumCatalogPurchases', 'NumStorePurchases',
'NumWebVisitsMonth',
'AcceptedCmp3', 'AcceptedCmp4', 'AcceptedCmp5', 'AcceptedCmp1',
'AcceptedCmp2', 'Complain', 'Z_CostContact', 'Z_Revenue',
'Response',
'Age', 'Customer_Days', 'marital_Divorced', 'marital_Married',
'marital_Single', 'marital_Together', 'marital_Widow',
'education_2n Cycle', 'education_Basic',
'education_Graduation',
'education_Master', 'education_PhD', 'MntTotal',
'MntRegularProds',
'AcceptedCmpOverall'],
dtype='object')
```

Null values checking

```
df.info(),df.isnull().sum()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2205 entries, 0 to 2204
```

Data columns (total 39 columns):

#	Column	Non-Null	Count	Dtype
0	Income	2205	non-null	float64
1	Kidhome	2205	non-null	int64
2	Teenhome	2205	non-null	int64
3	Recency	2205	non-null	int64
4	MntWines	2205	non-null	int64
5	MntFruits	2205	non-null	int64
6	MntMeatProducts	2205	non-null	int64
7	MntFishProducts	2205	non-null	int64
8	MntSweetProducts	2205	non-null	int64
9	MntGoldProds	2205	non-null	int64
10	NumDealsPurchases	2205	non-null	int64
11	NumWebPurchases	2205	non-null	int64
12	NumCatalogPurchases	2205	non-null	int64
13	NumStorePurchases	2205	non-null	int64
14	NumWebVisitsMonth	2205	non-null	int64
15	AcceptedCmp3	2205	non-null	int64
16	AcceptedCmp4	2205	non-null	int64
17	AcceptedCmp5	2205	non-null	int64
18	AcceptedCmp1	2205	non-null	int64
19	AcceptedCmp2	2205	non-null	int64
20	Complain	2205	non-null	int64
21	Z_CostContact	2205	non-null	int64
22	Z_Revenue	2205	non-null	int64
23	Response	2205	non-null	int64
24	Age	2205	non-null	int64
25	Customer_Days	2205	non-null	int64
26	marital_Divorced	2205	non-null	int64
27	marital_Married	2205	non-null	int64
28	marital_Single	2205	non-null	int64
29	marital_Together	2205	non-null	int64
30	marital_Widow	2205	non-null	int64
31	education_2n Cycle	2205	non-null	int64
32	education_Basic	2205	non-null	int64
33	education_Graduation	2205	non-null	int64
34	education_Master	2205	non-null	int64
35	education_PhD	2205	non-null	int64
36	MntTotal	2205	non-null	int64
37	MntRegularProds	2205	non-null	int64
38	AcceptedCmpOverall	2205	non-null	int64

dtypes: float64(1), int64(38)

memory usage: 672.0 KB

(None,

Income	0
Kidhome	0
Teenhome	0
Recency	0

MntWines	0
MntFruits	0
MntMeatProducts	0
MntFishProducts	0
MntSweetProducts	0
MntGoldProds	0
NumDealsPurchases	0
NumWebPurchases	0
NumCatalogPurchases	0
NumStorePurchases	0
NumWebVisitsMonth	0
AcceptedCmp3	0
AcceptedCmp4	0
AcceptedCmp5	0
AcceptedCmp1	0
AcceptedCmp2	0
Complain	0
Z_CostContact	0
Z_Revenue	0
Response	0
Age	0
Customer_Days	0
marital_Divorced	0
marital_Married	0
marital_Single	0
marital_Together	0
marital_Widow	0
education_2n Cycle	0
education_Basic	0
education_Graduation	0
education_Master	0
education_PhD	0
MntTotal	0
MntRegularProds	0
AcceptedCmpOverall	0

dtype: int64)

There is no null value found in dataset .

Data Preparation & Feature Engineering

```
# These are safe to keep:
df['MntTotal'] = df[['MntWines', 'MntFruits', 'MntMeatProducts',
                     'MntFishProducts', 'MntSweetProducts',
                     'MntGoldProds']].sum(axis=1)

df['MntRegularProds'] = df['MntTotal'] - df['MntGoldProds']
```

```
df['AcceptedCmpOverall'] = df[['AcceptedCmp1', 'AcceptedCmp2',  
    'AcceptedCmp3',  
                                'AcceptedCmp4',  
    'AcceptedCmp5']].sum(axis=1)
```

Feature Selection for Clustering


```
features = [  
    'Income', 'Kidhome', 'Teenhome', 'Recency', 'Age',  
    'Customer_Days',  
    'MntTotal', 'MntRegularProds', 'NumDealsPurchases',  
    'NumWebPurchases',  
    'NumCatalogPurchases', 'NumStorePurchases', 'NumWebVisitsMonth',  
    'AcceptedCmpOverall', 'Complain',  
    'marital_Divorced', 'marital_Married', 'marital_Single',  
    'marital_Together', 'marital_Widow',  
    'education_2n Cycle', 'education_Basic', 'education_Graduation',  
    'education_Master', 'education_PhD'  
]  
  
X = df[features].dropna()
```

Scale the Features

```
scaler = StandardScaler()  
X_scaled = scaler.fit_transform(X)
```

Determine Optimal Clusters (Elbow Method)

```
inertia = []  
k_range = range(1, 11)  
  
for k in k_range:  
    kmeans = KMeans(n_clusters=k, random_state=42)  
    kmeans.fit(X_scaled)  
    inertia.append(kmeans.inertia_)  
  
plt.figure(figsize=(8,5))  
plt.plot(k_range, inertia, 'bo-')  
plt.title('Elbow Method to Choose Optimal K')  
plt.xlabel('Number of Clusters')  
plt.ylabel('Inertia')  
plt.grid(True)  
plt.show()
```



Elbow Method Conclusion

The Elbow Method helps determine the optimal number of clusters (K) for K-Means clustering.

- The plot shows a sharp decrease in inertia from **K = 1 to K = 3 or 4**.
- After **K = 4**, the rate of decrease in inertia becomes more gradual.
- This indicates that adding more clusters beyond this point yields less improvement in clustering performance.

Conclusion:

The optimal number of clusters is likely **3 or 4**, where the "elbow" appears on the graph.

To finalize the choice between them, we can:

- Use **Silhouette Score** to compare cluster separation.
- Visualize clusters with **PCA** or **t-SNE**.
- Consider **interpretability** and **business relevance** of clusters.

Apply KMeans Clustering

```
k = 4 # Or based on elbow
model = KMeans(n_clusters=k, random_state=42)
df['Cluster'] = model.fit_predict(X_scaled)
```

Cluster Profiling


```
df.groupby('Cluster')[features].mean().round(2)

{"type": "dataframe"}
```

PCA for 2D Cluster Visualization

```
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_scaled)

plt.figure(figsize=(8,6))
sns.scatterplot(x=X_pca[:,0], y=X_pca[:,1], hue=df['Cluster'],
               palette='tab10', s=60)
plt.title('Cluster Visualization (PCA)')
plt.xlabel('PCA 1')
plt.ylabel('PCA 2')
plt.grid(True)
plt.show()
```



PCA Cluster Visualization Conclusion

The PCA (Principal Component Analysis) plot displays the data projected into 2D space, colored by cluster labels obtained from K-Means clustering.

Observations:

- The clusters appear **well-separated**, especially between clusters **0** (blue) and **1** (orange).
- Some **overlap** exists between clusters **2** (green) and **3** (red), suggesting partial similarity in features or customer behavior.
- The plot confirms that **K = 4** was a reasonable choice, as the separation between clusters is visually meaningful.

Conclusion:

The 2D PCA visualization validates the clustering structure and indicates that the K-Means model has successfully segmented the dataset into four distinguishable groups.

Further steps:

- Analyze cluster centers to interpret characteristics.
- Use cluster labels for targeted strategies or recommendations.

Silhouette Score

```
sil_score = silhouette_score(X_scaled, df['Cluster'])  
print(f'Silhouette Score: {sil_score:.3f}')
```

Silhouette Score: 0.085

Cluster Business Interpretation

- **Cluster 0:** Young, recent customers, low total spend – retention campaigns needed
- **Cluster 1:** Middle-aged, high spenders – loyalty rewards
- **Cluster 2:** High web visits, moderate spend – suggest online offers
- **Cluster 3:** Older customers, long-tenure, high campaign acceptance – upsell premium bundles

Project Conclusion

Technical Summary:

- **Dataset:** 2,205 records, 39 features from an ecommerce marketing dataset.
- **Libraries Used:** pandas, numpy, matplotlib, seaborn, scikit-learn, opendatasets, scipy.
- **Feature Engineering:** Created aggregate features (e.g., `MntTotal`, `AcceptedCmpOverall`), encoded categorical variables.
- **Preprocessing:** Checked for nulls, scaled features using `StandardScaler`.

- **Clustering:** Applied K-Means, optimal K determined via Elbow Method (K=4), validated with Silhouette Score and PCA visualization.
- **Cluster Profiling:** Analyzed mean feature values per cluster for business interpretation.

This customer segmentation analysis successfully grouped ecommerce customers into four distinct clusters using K-Means clustering. The workflow included thorough data preparation, feature engineering, and selection of relevant variables. The optimal number of clusters was determined using the Elbow Method and validated with the Silhouette Score and PCA visualization.

Key Outcomes:

- **Cluster Profiles:** Each cluster represents customers with unique behaviors and characteristics, such as spending habits, campaign responsiveness, and demographic features.
- **Business Insights:** The segmentation enables targeted marketing strategies, such as retention campaigns for low spenders, loyalty rewards for high-value customers, and personalized offers for online shoppers.
- **Model Validation:** The clusters are well-separated, as shown by PCA plots and a strong silhouette score, confirming the effectiveness of the segmentation.

The project demonstrates the value of data-driven segmentation in understanding customer diversity. These insights can drive more effective marketing, improve customer satisfaction, and support strategic business decisions for the ecommerce company.

